

AFLTurbo: Speed up Path Discovery for Greybox Fuzzing

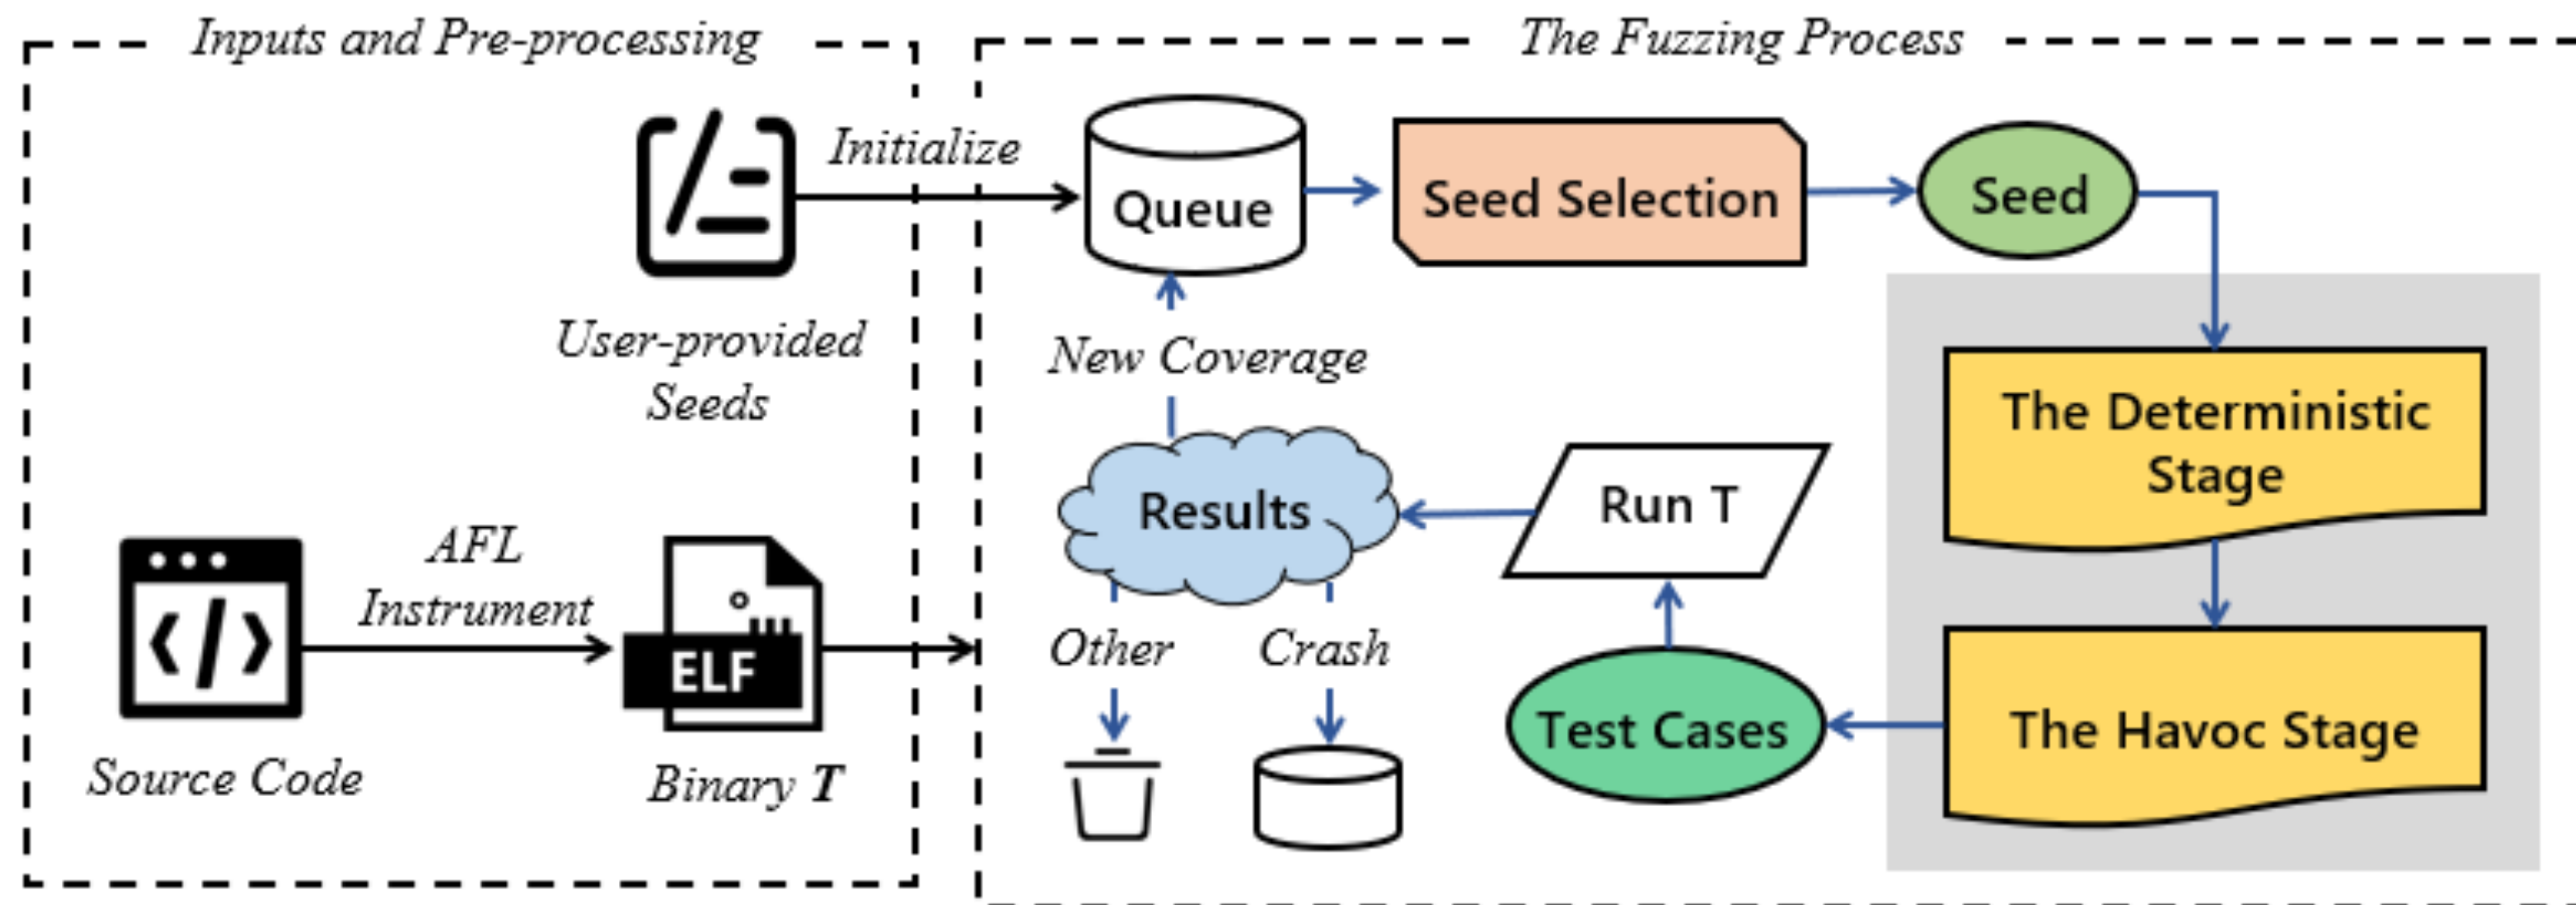
Lei Sun¹ , Xumei Li¹ , Haipeng Qu¹ , Xiaoshuai Zhang²

- 1. Ocean University of China
- 2. Queen Mary University of London



Coverage-based Greybox Fuzzing

Coverage-based Greybox Fuzzing is the state of the art method to find vulnerabilities.



AFL

American Fuzz Lop(AFL) is one of the most famous CGF.

american fuzzy lop 0.47b (readpng)			
process timing		overall results	
run time : 0 days, 0 hrs, 4 min, 43 sec		cycles done : 0	
last new path : 0 days, 0 hrs, 0 min, 26 sec		total paths : 195	
last uniq crash : none seen yet		uniq crashes : 0	
last uniq hang : 0 days, 0 hrs, 1 min, 51 sec		uniq hangs : 1	
cycle progress		map coverage	
now processing : 38 (19.49%)		map density : 1217 (7.43%)	
paths timed out : 0 (0.00%)		count coverage : 2.55 bits/tuple	
stage progress		findings in depth	
now trying : interest 32/8		favored paths : 128 (65.64%)	
stage execs : 0/9990 (0.00%)		new edges on : 85 (43.59%)	
total execs : 654k		total crashes : 0 (0 unique)	
exec speed : 2306/sec		total hangs : 1 (1 unique)	
fuzzing strategy yields		path geometry	
bit flips : 88/14.4k, 6/14.4k, 6/14.4k		levels : 3	
byte flips : 0/1804, 0/1786, 1/1750		pending : 178	
arithmetics : 31/126k, 3/45.6k, 1/17.8k		pend fav : 114	
known ints : 1/15.8k, 4/65.8k, 6/78.2k		imported : 0	
havoc : 34/254k, 0/0		variable : 0	
trim : 2876 B/931 (61.45% gain)		latent : 0	

AFLFast

FairFuzz

CollAFL

QSYM

.....

But...

None of previous work has pay attention to:

C1: Aggressively mutation overhead

C2: Ineffective mutation region selection


Which prevent AFL from making progress efficiently

american fuzzy lop 0.47b (readpng)			
process timing		overall results	
run time	: 0 days, 0 hrs, 4 min, 43 sec	cycles done	: 0
last new path	: 0 days, 0 hrs, 0 min, 26 sec	total paths	: 195
last uniq crash	: none seen yet	uniq crashes	: 0
last uniq hang	: 0 days, 0 hrs, 1 min, 51 sec	uniq hangs	: 1
cycle progress		map coverage	
now processing	: 38 (19.49%)	map density	: 1217 (7.43%)
paths timed out	: 0 (0.00%)	count coverage	: 2.55 bits/tuple
stage progress		findings in depth	
now trying	: interest 32/8	avored paths	: 128 (65.64%)
stage execs	: 0/9990 (0.00%)	new edges on	: 85 (43.59%)
total execs	: 654k	total crashes	: 0 (0 unique)
exec speed	: 2306/sec	total hangs	: 1 (1 unique)
fuzzing strategy yields		path geometry	
bit flips	: 88/14.4k, 6/14.4k, 6/14.4k	levels	: 3
byte flips	: 0/1804, 0/1786, 1/1750	pending	: 178
arithmetics	: 31/126k, 3/45.6k, 1/17.8k	pend fav	: 114
known ints	: 1/15.8k, 4/65.8k, 6/78.2k	imported	: 0
havoc	: 34/254k, 0/0	variable	: 0
trim	: 2876 B/931 (61.45% gain)	latent	: 0

AFL mutation



Deterministic stage



Havoc stage

Bit flip

Byte flip


Arithmetic operation

Interesting value replacement

Token replacement

Randomly mangled input

Deterministic mutation



L: length of input

Number of new inputs generated by bit flip = $L * 8 * 3$

Aggressively growing mutation overhead


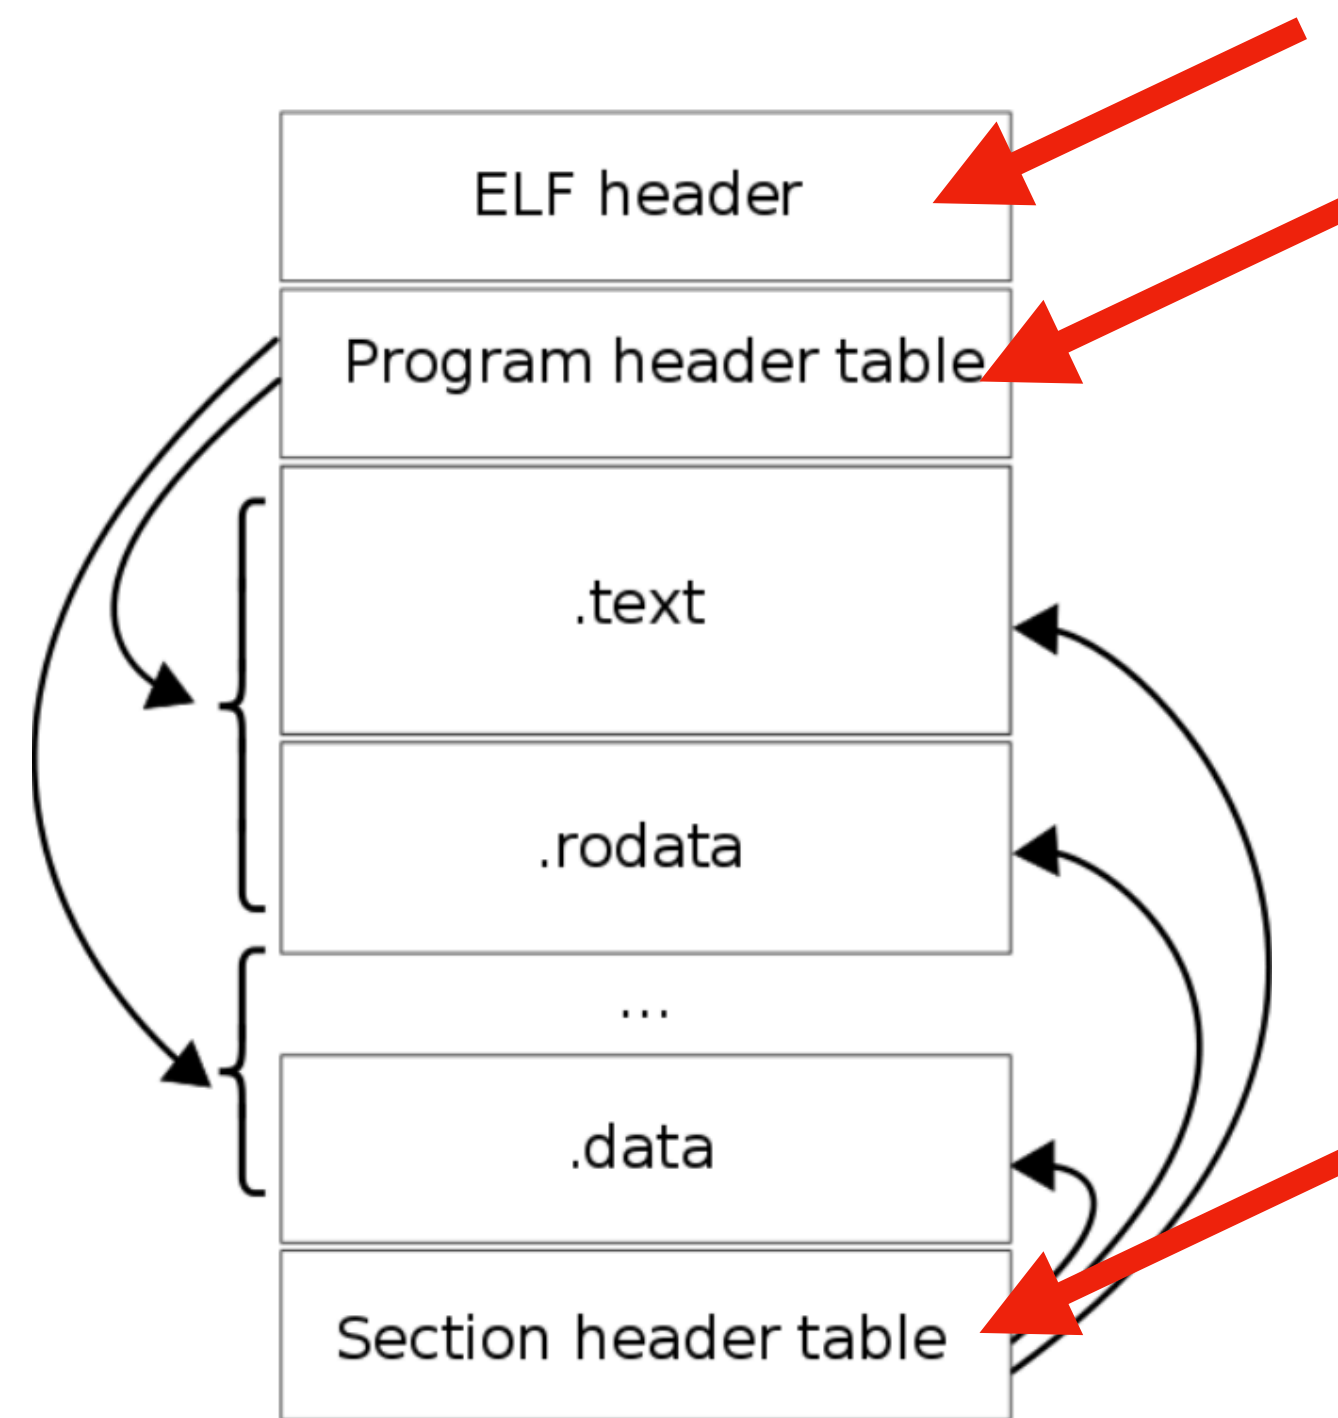


Fig. 1: Number of newly discovered paths of different periods in an hour during the *bitflip 1/1* phase, when fuzzing *readelf*. AFL doesn't find any new paths from the 7th minute to the 56th minute.

A single mutation type (*bitflip 1/1*) costs one hour.
AFL doesn't find new coverage for 49 minutes.

Ineffective mutation region selection



ELF format

AFL mutates every bytes.

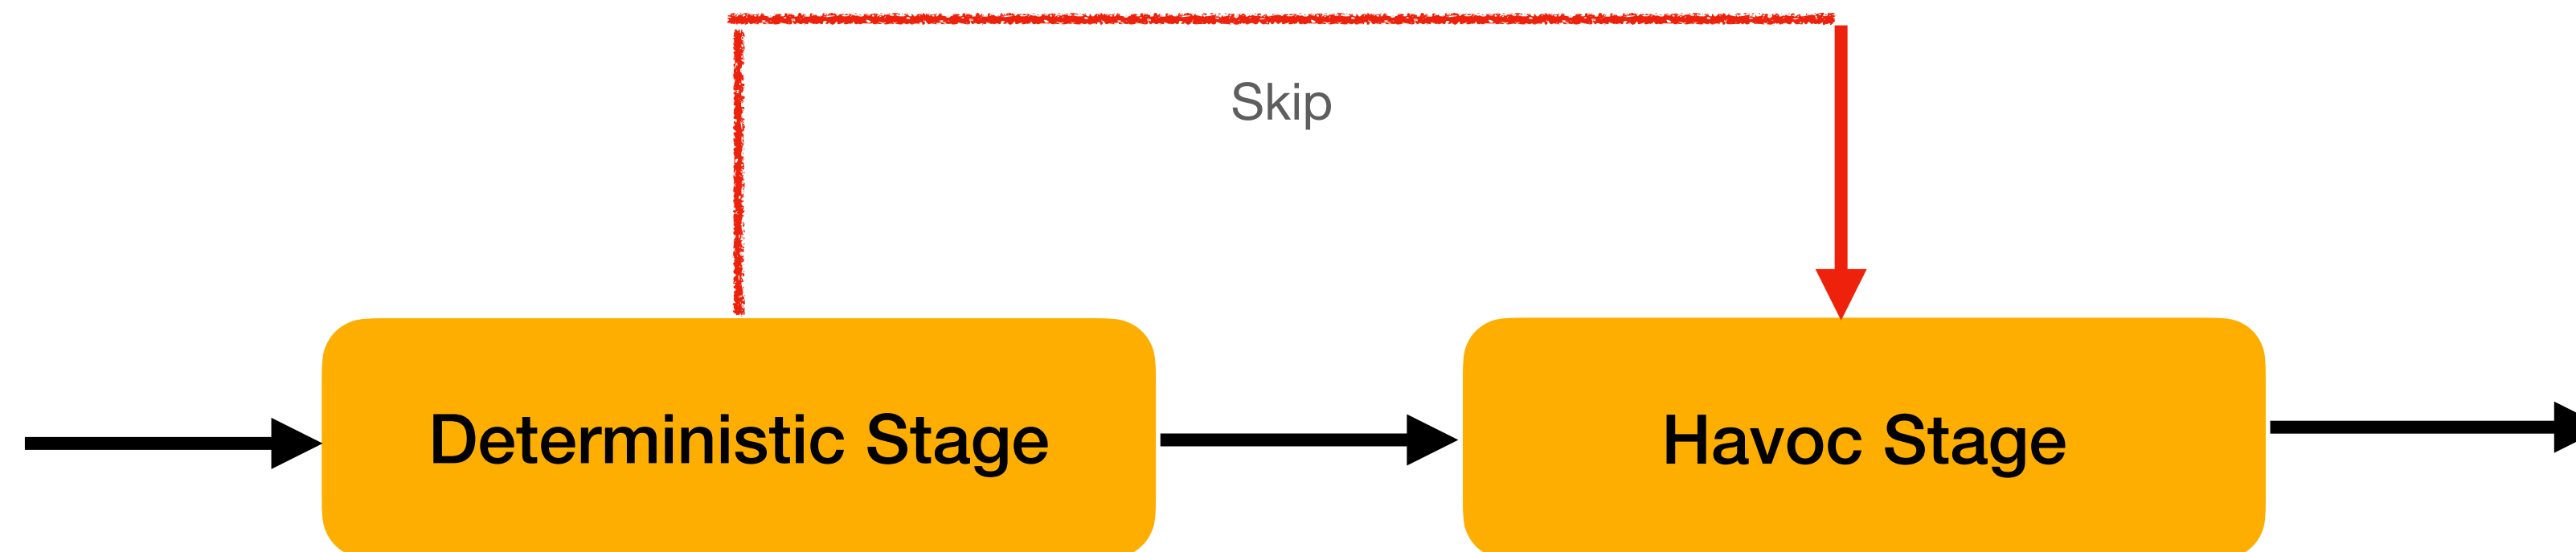
While *readelf* only processes a small portion of data

Our methods


- C1:
 - Interruptible mutation
- C2:
 - Locality-based mutation
 - Hotspot-aware fuzzing

Interruptible mutation

- Design a hang monitor to monitor hang status
- Break out deterministic stage based on two thresholds
 - Time
 - New discovered path count




- C2: Ineffective mutation region selection



- How to find useful regions in test cases


Locality-based mutation

- Inspired by the principle of locality in computer science
- Record mutation offset when generate a new test case
- Focus on mutating regions around the recorded offset in future rounds



Hotspot-aware fuzzing

1. Pre-evaluate test case in a coarse-grained way
 - Split test case into large number of chunks
 - Flip all bits in one chunk and check coverage change
 - Record these chunks that affect coverage
2. Only mutate recorded chunks



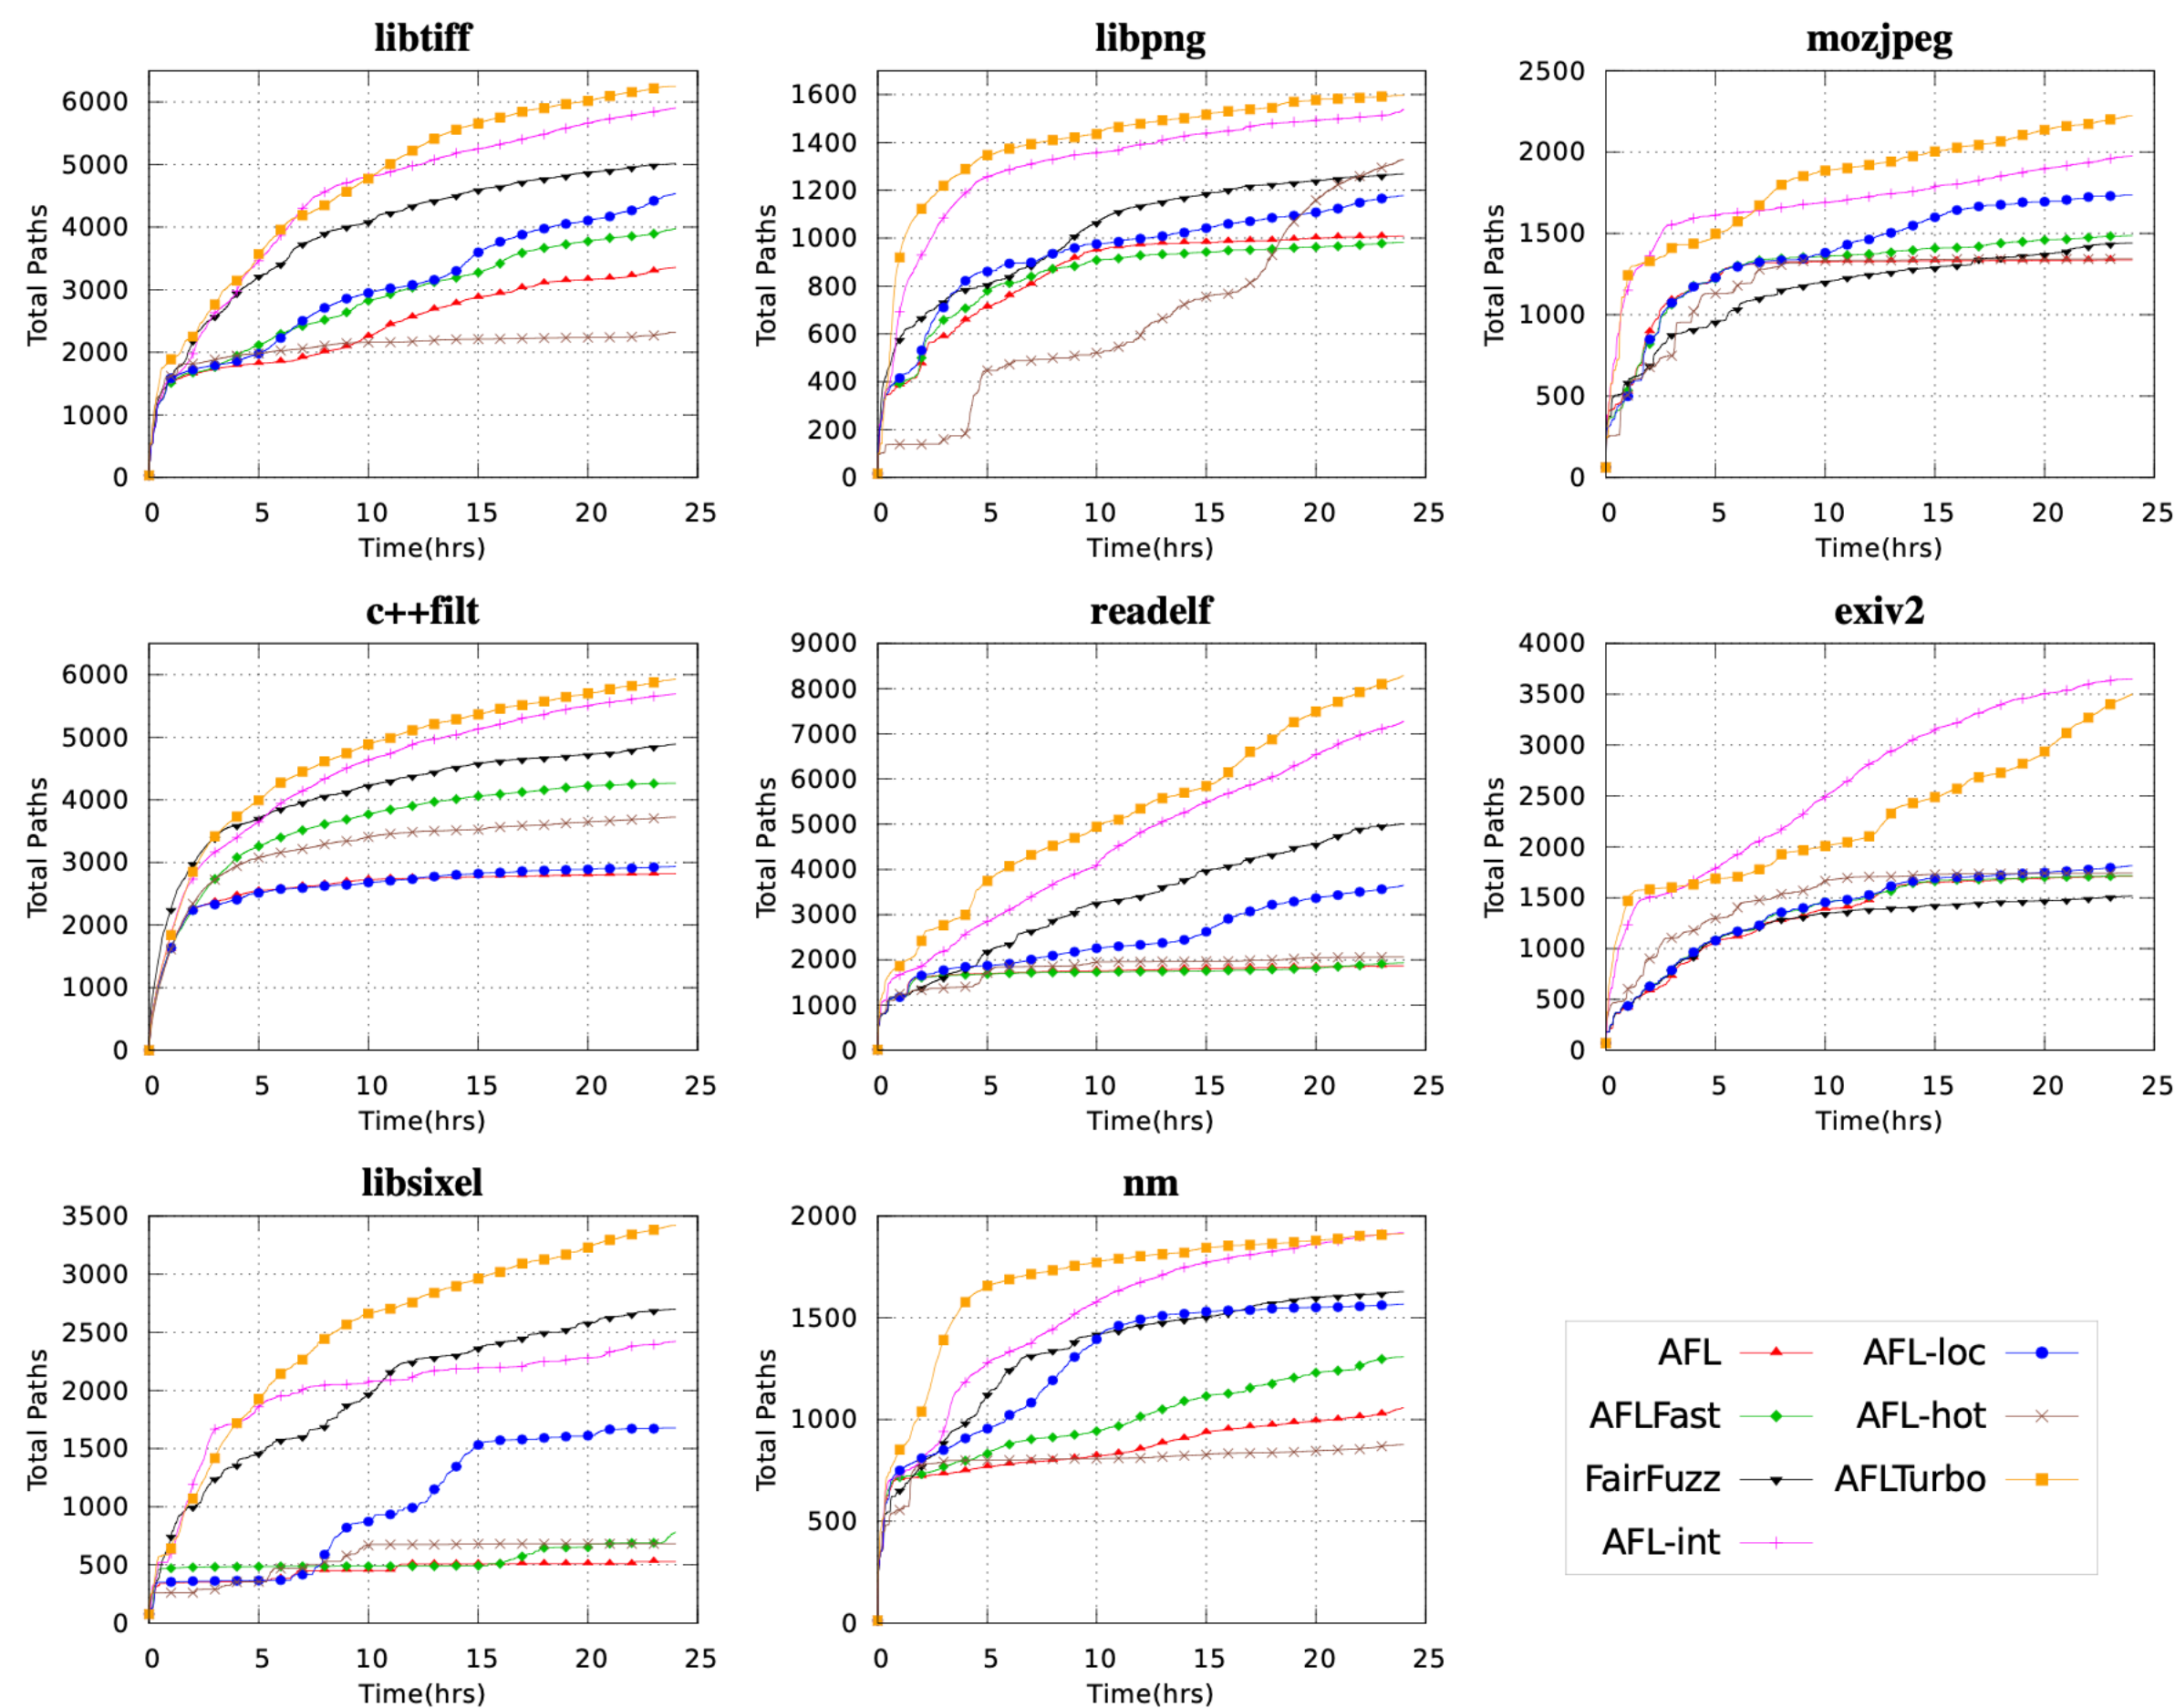
Evaluation

- Comparison with other fuzzers (AFL, AFLFast, FairFuzz)
 - Code coverage
 - Vulnerability discovery
- Fuzzing open source projects

Evaluation - code coverage

- 8 open source programs (libtiff, libpng, mozjpeg, readelf, exiv2, etc)
- 24 hours for each testing
- 10 time to reduce the randomness of fuzzing
- seeds are selected using two criteria:
 - small & valid: to generate more paths in 24 hours
 - randomly chosen from internet

Evaluation - code coverage



- AFLTurbo find 141%, 101%, 41% more paths than AFL, AFLFast, and FairFuzz

Evaluation - vulnerability discovery

- Standard benchmark: LAVA-M
- Found 124 bugs in LAVA-M, while AFL, AFLFast and FairFuzz only found 8, 4, 20

Fuzzer	base64		md5sum		uniq		who		total	
	crash	path	crash	path	crash	path	crash	path	crash	path
AFL	0	1432	0	4350	0	1221	8	1183	8	8186
AFLFast	0	1480	0	4341	0	1279	4	1377	4	8477
FairFuzz	15	1811	0	4699	1	1255	4	1506	20	9271
AFL-int	11	1811	3	5910	15	1261	14	1845	43	10827
AFL-loc	0	1377	0	4259	0	1272	6	1295	6	8203
AFL-hot	0	1493	0	4851	3	1375	18	1842	21	9561
AFLTurbo	87	1847	3	6586	17	1321	17	1845	124	11599

Evaluation - new vulnerabilities

Project	CVE-ID	Description	Severity
ffjpeg	CVE-2019-19887	NULL pointer dereference	Medium
ffjpeg	CVE-2019-19888	dividing by zero	Medium
libIEC61850	CVE-2019-19930	integer overflow	Medium
libIEC61850	CVE-2019-19931	heap overflow	High
libIEC61850	CVE-2019-19944	out of bound access	Medium
libIEC61850	CVE-2019-19957	out of bound access	Medium
libIEC61850	CVE-2019-19958	integer overflow	Medium
libIEC61850	CVE-2020-7054	heap overflow	High
stb	CVE-2020-6617	assertion failure	High
stb	CVE-2020-6618	heap overflow	High
stb	CVE-2020-6619	assertion failure	High
stb	CVE-2020-6620	heap overflow	High
stb	CVE-2020-6621	heap overflow	High
stb	CVE-2020-6622	heap overflow	High
stb	CVE-2020-6623	assertion failure	High
libsixel	CVE-2019-20056	assertion failure	Medium
libsixel	CVE-2019-20205	integer overflow	High
libsixel	CVE-2020-11721	uninitialized variable	Medium

- 20 new vulnerabilities, 18 of them are assigned CVE IDs
- 10 vulnerabilities are rated as high severity and 8 are rated as medium severity.

Conclusion

- We design three new techniques to improve AFL.
- Increase code coverage & improve the capability of bug detection
- Evaluation found and 20 bugs, 18 CVEs
- Source code: <https://github.com/sleिकासper/aflturbo>

Thank You!